

Correlations Between Complexity Measures

Chris Bentz

January 7, 2021

Load Libraries

If the libraries are not installed yet, you need to install them using, for example, the command: `install.packages("ggplot2")`.

```
library(readr)
library(ggplot2)
library(gridExtra)
library(GGally)
library(Hmisc)
library(ggrepel)
```

Load the Data

The participants' results are loaded as csv files directly from the github repository into separate data frames. We only use the name of the first author (lower case) to name the data frame.

```
#Track A (Parallel Bible Corpus, PBC)
gutierrez.results <- read_csv("https://raw.githubusercontent.com/IWMLC/language-complexity-metrics/master/PBCt.csv")
# remove the parentheses in column names
colnames(gutierrez.results) <- sub("\\(", "", colnames(gutierrez.results))
colnames(gutierrez.results) <- sub("\\)", "", colnames(gutierrez.results))
# replace "+" by "."
colnames(gutierrez.results) <- gsub("\\+", ".", colnames(gutierrez.results))
oh.results <- read_csv("https://raw.githubusercontent.com/IWMLC/language-complexity-metrics/master/PBCt.csv")

#TRACK B (Universal Dependencies, UD)
brunato.results <- read_csv("https://raw.githubusercontent.com/IWMLC/language-complexity-metrics/master/UDbrunato.csv")
coltekin.results <- read_csv("https://raw.githubusercontent.com/IWMLC/language-complexity-metrics/master/UDcoltekin.csv")
semenuks.results <- read_csv("https://raw.githubusercontent.com/IWMLC/language-complexity-metrics/master/UDsemenuks.csv")
sinemaki.results <- read_csv("https://raw.githubusercontent.com/IWMLC/language-complexity-metrics/master/UDsinemaki.csv")
sozinova.results <- read_csv("https://raw.githubusercontent.com/IWMLC/language-complexity-metrics/master/UDsozinova.csv")
```

Sanity check, look at the number of rows and columns of the data frames.

```
#Track A (should be 49 rows)
track.a.rows <- c(nrow(gutierrez.results), nrow(oh.results))
print(track.a.rows) # this corresponds to the number of languages
```

```
## [1] 49 49
```

```

track.a.cols <- c(ncol(gutierrez.results), ncol(oh.results))
print(track.a.cols) # this is the number of measures per team

## [1] 14 5

#Track B (should be 63 rows)
track.b.rows <- c(nrow(brunato.results), nrow(coltekin.results), nrow(semenuks.results), nrow(sinnemaki.results))
print(track.b.rows) # this corresponds to the number of languages

## [1] 63 63 63 63 63

track.b.cols <- c(ncol(brunato.results), ncol(coltekin.results), ncol(semenuks.results), ncol(sinnemaki.results))
print(track.b.cols) # this is the number of measures per team

## [1] 13 8 4 8 4

```

Preprocessing

Put data into a single data frame.

```

track.a <- cbind(gutierrez.results, oh.results[, 3:ncol(oh.results)])
track.b <- cbind(brunato.results, coltekin.results[, 3:ncol(coltekin.results)],
                 semenuks.results[, 3:ncol(semenuks.results)],
                 sinnemaki.results[, 3:ncol(sinnemaki.results)],
                 sozinova.results[, 3:ncol(sozinova.results)])

```

Check data frames by looking at the first six rows.

```

head(track.a)

##      id      language GM_H1gram GM_H3gram      GM_TTR GM_TTR.H1 GM_TTR.H3
## 1 aey      Amele 0.6119331 0.780252 0.06065136 0.04000000 0.03947368
## 2 amp      Alamlak 0.7057048 0.782454 0.10978104 0.05000000 0.06382979
## 3 ape      Bukiyip 0.6311849 0.763648 0.05561431 0.03333333 0.03614458
## 4 apu      Apurinã 0.5809332 0.567297 0.09235873 0.02985075 0.02857143
## 5 arn      Mapudungun 0.5744377 0.780808 0.08595837 0.04347826 0.03488372
## 6 arz      Egyptian Arabic 0.8083907 0.887509 0.12832227 0.18181818 0.25000000
##      GM_TTR.H1.H3 GM_H1gram_fullyparallelised GM_H3gram_fullyparallelised
## 1 0.04081633      0.5684752      0.5908496
## 2 0.09090909      0.6732652      0.6431306
## 3 0.04081633      0.6519671      0.5913951
## 4 0.03571429      0.5924908      0.5236397
## 5 0.03333333      0.5983614      0.5964545
## 6 0.22222222      0.7258967      0.7482857
##      GM_TTR_fullyparallelised GM_TTR.H1_fullyparallelised
## 1 0.1342350      0.03125000
## 2 0.2033711      0.07692308
## 3 0.1193411      0.04166667
## 4 0.2053489      0.04651163
## 5 0.1457532      0.04166667
## 6 0.3107240      0.22222222
##      GM_TTR.H3_fullyparallelised GM_TTR.H1.H3_fullyparallelised      O_MC
## 1 0.03636364      0.03225806 -0.5125000
## 2 0.06060606      0.06818182 -0.4288462
## 3 0.03333333      0.03947368 -0.5021739

```

```
## 4          0.03508772          0.03488372 -0.5660714
## 5          0.04255319          0.04000000 -0.5351852
## 6          0.25000000          0.23076923 -0.4086207
##      O_WID O_SID
## 1 1.208351    NA
## 2 1.161317    NA
## 3 0.927140    NA
## 4 1.656120    NA
## 5 1.226694    NA
## 6 1.972695    NA
```

```
head(track.b)
```

```
##      id language BV_n_tokens BV_char_per_tok BV_verbal_head_per_sent
## 1   afr Afrikaans  0.6915035      0.7406356      0.5973389
## 2   ara  Arabic   1.0000000      0.5700036      0.7373122
## 3   bul Bulgarian 0.3804936      0.7071177      0.4128931
## 4   cat  Catalan  0.8656840      0.6355104      0.6404372
## 5   ces_cac Czech  0.5430299      0.7629205      0.4396123
## 6 ces_fictree Czech 0.3553259      0.6460943      0.5242218
## BV_verbal_root_perc BV_avg_token_per_clause BV_avg_links_len BV_avg_max_depth
## 1      0.7934703      0.8065034      0.9435208      0.7908597
## 2      0.2567031      0.9005952      0.7894060      1.0000000
## 3      0.8396519      0.5572851      0.6627371      0.4848930
## 4      0.8565512      1.0000000      0.8081980      0.7408015
## 5      0.7177322      0.7180483      0.7227964      0.6323912
## 6      0.8266025      0.4074251      0.6584553      0.4199076
## BV_avg_subordinate_chain_len BV_subordinate_pre BV_subordinate_post
## 1      0.6690569      0.195677426      0.7325159
## 2      0.8753356      0.006573444      1.0000000
## 3      0.4185603      0.140466212      0.4342773
## 4      0.6893540      0.147016174      0.7545315
## 5      0.4598093      0.092272362      0.5474183
## 6      0.4772712      0.130639721      0.5117906
## BV_avg_verb_edges CR_inflection_accuracy CR_ttr CR_msp CR_mfe
## 1      0.8840035      0.6670443 0.1853184 1.219609 6.427348
## 2      0.6585259      0.1000941 0.2988734 1.475221 6.976586
## 3      0.6340592      0.3490676 0.3770606 1.514225 9.336138
## 4      0.8522001      0.5170170 0.2757109 1.266932 6.275705
## 5      0.6145715      0.3340556 0.4492450 1.588910 11.145461
## 6      0.6190423      0.3130136 0.3712954 1.633426 11.673832
## CR_cfe_form_feat CR_cfe_feat_form S_idMean S_idSD SI_dm SI_hm
## 1 7.53990e-06 8.95090e-06 1.442039 1.0039941 1.0000000 0
## 2 2.19027e-05 6.96940e-06 1.407543 1.0389387 1.0000000 0
## 3 4.78590e-06 1.11903e-05 1.554631 1.0983221 1.0000000 0
## 4 1.17581e-05 1.05572e-05 1.567901 1.1048304 1.0000000 0
## 5 1.37917e-05 6.86860e-06 1.754058 1.0347654 0.8430774 0
## 6 1.55421e-05 1.07678e-05 1.821586 0.9678328 0.9334056 0
## SI_dep_dl SI_double_dl SI_head_dl SI_zero_dl SBS_INF SBS_DER
## 1 2.346654 NA NA NA 0.199 0.093
## 2 1.830569 NA NA NA 0.955 0.032
## 3 2.275108 NA NA NA 1.020 0.209
## 4 2.641555 NA NA NA 0.845 0.059
## 5 1.745720 NA NA 1.837112 1.439 0.117
## 6 1.425285 NA NA 1.609121 1.461 0.148
```


Plot Correlations by Track

TRACK A

We here plot correlations between selected measures of the respective track (trying to exclude the ones which are somewhat redundant). The Spearman correlation coefficient is reported instead of the Pearson correlation coefficient. This is because we are only interested whether there is a correlation between the rankings of complexities, regardless of whether this is a linear relationship. We therefore also use the local regression smoothers in the plots (loess) rather than linear models (lm). Note: warning messages are disabled here as there are datasets with NAs, and for each plot this throws a warning message using the ggpairs() plotting function. NAs are dealt with by removing the entire row, containing an NA value.

```
#remove the first two columns for plotting
track.a.short <- track.a[, 3:ncol(track.a)]

track.a.plot <- ggpairs(track.a.short,
                        lower = list(continuous = wrap("smooth_loess", alpha = 0.3,
                                                         lwd = 0.5, size = 2)),
                        upper = list(continuous = wrap('cor', method = "spearman")) +
                        theme_bw() +
                        theme(axis.text.x = element_text(angle = 90, hjust = 1))
print(track.a.plot)
```



Safe plot to file.

```
ggsave("~/Github/ComplexityMetaAnalyses/Figures/Corrs/TrackA/track_a_plot.pdf", track.a.plot,
        dpi = 300, scale = 1, width = 15, height = 15, device = cairo_pdf)
```

TRACK B

Same for the Track B data. Not all measures are included here (there would be 27). To include them all, the “columns” argument in the code below might be removed.

```
#remove the first to columns for plotting
track.b.short <- track.b[, 3:ncol(track.b)]

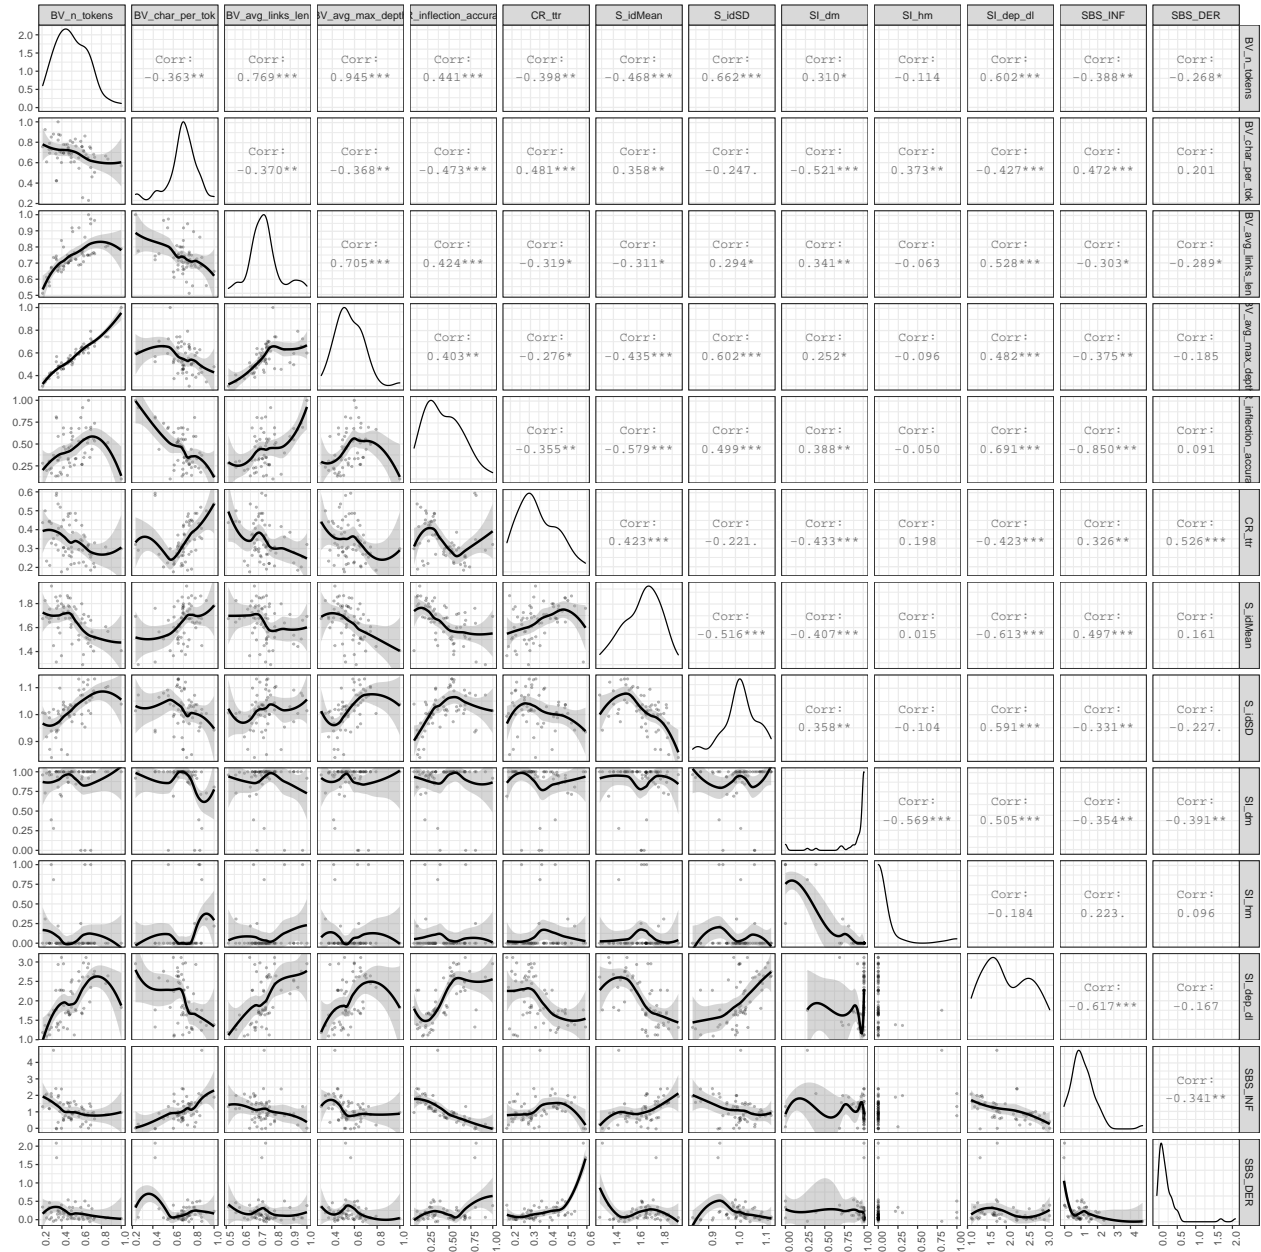
track.b.plot <- ggpairs(track.b.short, progress = TRUE,
```

```

lower = list(continuous = wrap("smooth_loess", alpha = 0.3,
                                lwd = 0.5, size = 2)),
upper = list(continuous = wrap('cor', method = "spearman")),
columns = c("BV_n_tokens", "BV_char_per_tok", "BV_avg_links_len",
            "BV_avg_max_depth", "CR_inflection_accuracy", "CR_ttr",
            "S_idMean", "S_idSD", "SI_dm", "SI_hm", "SI_dep_dl",
            "SBS_INF", "SBS_DER")) +
theme_bw() +
theme(axis.text.x = element_text(angle = 90, hjust = 1))

```

```
print(track.b.plot)
```



Save plot to file.

```
ggsave("~/Github/ComplexityMetaAnalyses/Figures/Corrs/TrackB/track_b_plot.pdf", track.b.plot,
        dpi = 300, scale = 1, width = 15, height = 15, device = cairo_pdf)
```

Significant Correlations after Bonferroni Correction

Not all of the correlations displayed above are going to be significant. We only select the ones still significant after correcting for multiple testing. Therefore, first calculate Spearman correlations and uncorrected p-values using the function `rcorr()`. These are then stored in a data frame where the first two columns give the names of the correlated measures.

TRACK A

```
#transform the data frame to a matrix
track.a.matrix <- as.matrix(track.a.short)
#apply the rcorr function to this matrix to get matrices of Spearman correlations and uncorrected p-values
track.a.cor <- rcorr(track.a.matrix, type = "spearman")$r
track.a.pvalues <- rcorr(track.a.matrix, type = "spearman")$P
track.a.n <- rcorr(track.a.matrix, type = "spearman")$n
#convert these matrices to a data frame again
track.a.df <- data.frame(row = rownames(track.a.pvalues)[upper.tri(track.a.pvalues)],
                        col = colnames(track.a.pvalues)[upper.tri(track.a.pvalues)]],
                        pvalue = track.a.pvalues[upper.tri(track.a.pvalues)],
                        corr = track.a.cor[upper.tri(track.a.cor)],
                        num = track.a.n[upper.tri(track.a.n)])
head(track.a.df)
```

##	row	col	pvalue	corr	num
## 1	GM_H1gram	GM_H3gram	4.549026e-03	0.3986735	49
## 2	GM_H1gram	GM_TTR	1.317108e-02	0.3518367	49
## 3	GM_H3gram	GM_TTR	3.502930e-01	0.1363265	49
## 4	GM_H1gram	GM_TTR.H1	2.477667e-04	0.5007655	49
## 5	GM_H3gram	GM_TTR.H1	3.771262e-09	0.7252501	49
## 6	GM_TTR	GM_TTR.H1	7.641920e-10	0.7460706	49

TRACK B

Same as above for Track A.

```
track.b.matrix <- as.matrix(track.b.short)
track.b.cor <- rcorr(track.b.matrix, type = "spearman")$r
```

```
## Warning in sqrt(npair - 2): NaNs produced
```

```
track.b.pvalues <- rcorr(track.b.matrix, type = "spearman")$P
```

```
## Warning in sqrt(npair - 2): NaNs produced
```

```
track.b.n <- rcorr(track.b.matrix, type = "spearman")$n
```

```
## Warning in sqrt(npair - 2): NaNs produced
```



```
track.b.df <- data.frame(row = rownames(track.b.pvalues)[row(track.b.pvalues)[upper.tri(track.b.pvalues)],
col = colnames(track.b.pvalues)[col(track.b.pvalues)[upper.tri(track.b.pvalues)]],
pvalue = track.b.pvalues[upper.tri(track.b.pvalues)],
corr = track.b.cor[upper.tri(track.b.cor)],
num = track.b.n[upper.tri(track.b.n)])
#head(track.b.df)
```

Apply Bonferroni Correction

Apply the so-called Bonferroni correction, namely, multiply the p-values by the overall number of tests done. Arguably, this is the simplest, and also most conservative method for correcting the p-values. There are less-conservative alternatives such as the Holm-Bonferroni correction. Since the approach here is purely exploratory, and we have many measures and hence pairwise correlations anyways, we decided to go for the most conservative method.

```
# compute the overall number of tests, i.e. multiply the number of measures in each track with the same
n.test <- ncol(track.a.short)*(ncol(track.a.short) - 1) + ncol(track.b.short)*(ncol(track.b.short) - 1)
# add corrected pvalues to data frames
track.a.df$pvalue.correct <- track.a.df$pvalue*n.test
track.b.df$pvalue.correct <- track.b.df$pvalue*n.test
```

Remove all correlations which are not significant anymore. And then order them from highest to lowest coefficient.

```
# Track A
track.a.df <- track.a.df[track.a.df$pvalue.correct < 0.05, ]
track.a.df <- track.a.df[order(-track.a.df$corr), ]

# Track B
track.b.df <- track.b.df[track.b.df$pvalue.correct < 0.05, ]
track.b.df <- track.b.df[order(-track.b.df$corr), ]
```

Correlations still significant after Bonferroni correction for Track A:

```
print(track.a.df)
```

##	row	col	pvalue
## 31	GM_TTR	GM_TTR_fullyparallelised	0.000000e+00
## 10	GM_TTR.H1	GM_TTR.H3	0.000000e+00
## 65	GM_TTR.H1_fullyparallelised	GM_TTR.H1.H3_fullyparallelised	3.552714e-15
## 66	GM_TTR.H3_fullyparallelised	GM_TTR.H1.H3_fullyparallelised	6.217249e-15
## 60	GM_TTR.H3	GM_TTR.H1.H3_fullyparallelised	1.731948e-14
## 15	GM_TTR.H3	GM_TTR.H1.H3	1.243450e-14
## 49	GM_TTR.H1	GM_TTR.H3_fullyparallelised	6.483702e-14
## 50	GM_TTR.H3	GM_TTR.H3_fullyparallelised	7.926992e-14
## 42	GM_TTR.H1.H3	GM_TTR.H1_fullyparallelised	5.682121e-13
## 43	GM_H1gram_fullyparallelised	GM_TTR.H1_fullyparallelised	6.787904e-13
## 13	GM_TTR	GM_TTR.H1.H3	6.155076e-13
## 23	GM_H3gram	GM_H3gram_fullyparallelised	2.202238e-12
## 45	GM_TTR_fullyparallelised	GM_TTR.H1_fullyparallelised	2.940537e-12
## 34	GM_TTR.H1.H3	GM_TTR_fullyparallelised	3.945289e-12
## 62	GM_H1gram_fullyparallelised	GM_TTR.H1.H3_fullyparallelised	5.050182e-12
## 11	GM_H1gram	GM_TTR.H1.H3	1.898037e-12
## 87	GM_TTR_fullyparallelised	O_WID	5.985212e-12
## 7	GM_H1gram	GM_TTR.H3	3.868461e-12


```

## 41          GM_TTR.H3      GM_TTR.H1_fullyparallelised 5.331886e-10
## 59          GM_TTR.H1 GM_TTR.H1.H3_fullyparallelised 5.860898e-10
## 6           GM_TTR          GM_TTR.H1 7.641920e-10
## 39          GM_TTR      GM_TTR.H1_fullyparallelised 2.823907e-09
## 14          GM_TTR.H1          GM_TTR.H1.H3 2.268288e-09
## 81          GM_TTR          O_WID 5.209562e-09
## 5           GM_H3gram          GM_TTR.H1 3.771262e-09
## 32          GM_TTR.H1      GM_TTR_fullyparallelised 1.988542e-08
## 61          GM_TTR.H1.H3 GM_TTR.H1.H3_fullyparallelised 2.673641e-08
## 47          GM_H3gram      GM_TTR.H3_fullyparallelised 2.842440e-08
## 8           GM_H3gram          GM_TTR.H3 1.430519e-08
## 53 GM_H3gram_fullyparallelised GM_TTR.H3_fullyparallelised 6.724529e-08
## 16          GM_H1gram      GM_H1gram_fullyparallelised 1.151531e-07
## 40          GM_TTR.H1      GM_TTR.H1_fullyparallelised 1.242348e-07
## 55 GM_TTR.H1_fullyparallelised GM_TTR.H3_fullyparallelised 2.824460e-07
## 56          GM_H1gram GM_TTR.H1.H3_fullyparallelised 3.339067e-07
## 88 GM_TTR.H1_fullyparallelised O_WID 3.374804e-07
## 33          GM_TTR.H3      GM_TTR_fullyparallelised 4.658630e-07
## 37          GM_H1gram      GM_TTR.H1_fullyparallelised 5.418223e-07
## 54      GM_TTR_fullyparallelised GM_TTR.H3_fullyparallelised 6.189858e-07
## 51          GM_TTR.H1.H3      GM_TTR.H3_fullyparallelised 1.258620e-06
## 64      GM_TTR_fullyparallelised GM_TTR.H1.H3_fullyparallelised 1.425442e-06
## 9           GM_TTR          GM_TTR.H3 1.147100e-06
## 20          GM_TTR.H3      GM_H1gram_fullyparallelised 3.541727e-06
## 82          GM_TTR.H1          O_WID 5.154513e-06
## 57          GM_H3gram GM_TTR.H1.H3_fullyparallelised 1.419533e-05
## 63 GM_H3gram_fullyparallelised GM_TTR.H1.H3_fullyparallelised 1.515231e-05
## 21          GM_TTR.H1.H3      GM_H1gram_fullyparallelised 1.538187e-05
## 84          GM_TTR.H1.H3          O_WID 2.145134e-05
## 72          GM_TTR.H1.H3          O_MC 1.734576e-05
## 48          GM_TTR      GM_TTR.H3_fullyparallelised 4.770488e-05
## 89 GM_TTR.H3_fullyparallelised O_WID 5.298099e-05
##      corr num pvalue.correct
## 31 0.9293478 47 0.000000e+00
## 10 0.8964778 49 0.000000e+00
## 65 0.8667671 47 3.240075e-12
## 66 0.8629274 47 5.670131e-12
## 60 0.8560560 47 1.579537e-11
## 15 0.8493235 49 1.134026e-11
## 49 0.8466749 47 5.913137e-11
## 50 0.8451810 47 7.229417e-11
## 42 0.8297251 47 5.182095e-10
## 43 0.8282468 47 6.190568e-10
## 13 0.8193559 49 5.613430e-10
## 23 0.8181082 47 2.008441e-09
## 45 0.8155179 47 2.681769e-09
## 34 0.8128417 47 3.598103e-09
## 62 0.8105608 47 4.605766e-09
## 11 0.8095555 49 1.731010e-09
## 87 0.8089732 47 5.458514e-09
## 7 0.8030520 49 3.528037e-09
## 41 0.7610477 47 4.862680e-07
## 59 0.7598993 47 5.345139e-07
## 6 0.7460706 49 6.969431e-07

```

```
## 39 0.7398388 47 2.575403e-06
## 14 0.7320927 49 2.068679e-06
## 81 0.7314986 47 4.751120e-06
## 5 0.7252501 49 3.439391e-06
## 32 0.7121212 47 1.813550e-05
## 61 0.7076149 47 2.438361e-05
## 47 0.7066723 47 2.592306e-05
## 8 0.7062877 49 1.304633e-05
## 53 0.6930157 47 6.132771e-05
## 16 0.6840888 47 1.050196e-04
## 40 0.6828034 47 1.133021e-04
## 55 0.6684715 47 2.575907e-04
## 56 0.6654522 47 3.045229e-04
## 88 0.6652589 47 3.077821e-04
## 33 0.6593407 47 4.248671e-04
## 37 0.6565223 47 4.941420e-04
## 54 0.6540133 47 5.645151e-04
## 51 0.6402374 47 1.147862e-03
## 64 0.6377491 47 1.300003e-03
## 9 0.6314178 49 1.046155e-03
## 20 0.6188549 47 3.230055e-03
## 82 0.6106870 47 4.700916e-03
## 57 0.5874324 47 1.294615e-02
## 63 0.5858709 47 1.381891e-02
## 21 0.5855098 47 1.402826e-02
## 84 0.5774115 47 1.956362e-02
## 72 0.5723256 49 1.581934e-02
## 48 0.5570282 47 4.350685e-02
## 89 0.5542505 47 4.831866e-02
```

Correlations still significant after Bonferroni correction for Track B:

```
print(track.b.df)
```

```
##          row          col      pvalue
## 253      SI_dep_dl      SI_double_dl 0.000000e+00
## 276      SI_double_dl      SI_head_dl 0.000000e+00
## 300      SI_head_dl      SI_zero_dl 0.000000e+00
## 314      CR_msp          SBS_INF 0.000000e+00
## 16      BV_n_tokens      BV_avg_max_depth 0.000000e+00
## 7      BV_n_tokens      BV_avg_token_per_clause 0.000000e+00
## 20      BV_avg_token_per_clause      BV_avg_max_depth 0.000000e+00
## 24      BV_verbal_head_per_sent      BV_avg_subordinate_chain_len 0.000000e+00
## 105      CR_msp          CR_mfe 1.418865e-13
## 11      BV_n_tokens      BV_avg_links_len 1.880718e-13
## 315      CR_mfe          SBS_INF 3.888001e-13
## 21      BV_avg_links_len      BV_avg_max_depth 1.114653e-10
## 222      CR_inflection_accuracy      SI_dep_dl 1.905962e-09
## 15      BV_avg_token_per_clause      BV_avg_links_len 7.305712e-10
## 158      BV_avg_token_per_clause      S_idSD 9.133534e-10
## 50      BV_avg_token_per_clause      BV_avg_verb_edges 2.596476e-09
## 154      BV_n_tokens      S_idSD 3.423637e-09
## 37      BV_n_tokens      BV_subordinate_post 9.996910e-09
## 51      BV_avg_links_len      BV_avg_verb_edges 1.625949e-08
## 31      BV_verbal_head_per_sent      BV_subordinate_pre 1.992449e-08
```

## 36	BV_avg_subordinate_chain_len	BV_subordinate_pre	2.628625e-08
## 43	BV_avg_max_depth	BV_subordinate_post	3.838314e-08
## 151	CR_mfe	S_idMean	1.586801e-07
## 160	BV_avg_max_depth	S_idSD	1.748367e-07
## 211	BV_n_tokens	SI_dep_dl	4.647820e-07
## 46	BV_n_tokens	BV_avg_verb_edges	1.952684e-07
## 229	S_idSD	SI_dep_dl	8.127537e-07
## 215	BV_avg_token_per_clause	SI_dep_dl	1.004935e-05
## 41	BV_avg_token_per_clause	BV_subordinate_post	7.131573e-06
## 27	BV_avg_links_len	BV_avg_subordinate_chain_len	7.646974e-06
## 52	BV_avg_max_depth	BV_avg_verb_edges	7.820502e-06
## 216	BV_avg_links_len	SI_dep_dl	1.752547e-05
## 338	CR_ttr	SBS_DER	9.330432e-06
## 230	SI_dm	SI_dep_dl	4.464129e-05
## 165	CR_inflection_accuracy	S_idSD	3.690263e-05
## 318	S_idMean	SBS_INF	3.418259e-05
## 150	CR_msp	S_idMean	3.653193e-05
## 171	S_idMean	S_idSD	1.532149e-05
## 173	BV_char_per_tok	SI_dm	1.200458e-05
## 23	BV_char_per_tok	BV_avg_subordinate_chain_len	2.468442e-06
## 210	SI_dm	SI_hm	1.141080e-06
## 148	CR_inflection_accuracy	S_idMean	8.270105e-07
## 228	S_idMean	SI_dep_dl	2.411887e-07
## 322	SI_dep_dl	SBS_INF	1.948739e-07
## 224	CR_msp	SI_dep_dl	4.198780e-09
## 225	CR_mfe	SI_dep_dl	1.409184e-11
## 103	CR_inflection_accuracy	CR_mfe	2.220446e-16
## 312	CR_inflection_accuracy	SBS_INF	0.000000e+00
## 90	CR_inflection_accuracy	CR_msp	0.000000e+00
## NA	<NA>	<NA>	NA
##	corr num pvalue.correct		
## 253	1.0000000 3 0.000000e+00		
## 276	1.0000000 2 0.000000e+00		
## 300	1.0000000 2 0.000000e+00		
## 314	0.9475806 63 0.000000e+00		
## 16	0.9452765 63 0.000000e+00		
## 7	0.9151786 63 0.000000e+00		
## 20	0.8714478 63 0.000000e+00		
## 24	0.8623752 63 0.000000e+00		
## 105	0.7711899 63 1.294005e-10		
## 11	0.7687692 63 1.715215e-10		
## 315	0.7624043 63 3.545857e-10		
## 21	0.7050691 63 1.016563e-07		
## 222	0.6911624 58 1.738237e-06		
## 15	0.6824117 63 6.662809e-07		
## 158	0.6795795 63 8.329783e-07		
## 50	0.6658986 63 2.367986e-06		
## 154	0.6621544 63 3.122357e-06		
## 37	0.6471294 63 9.117181e-06		
## 51	0.6400250 63 1.482865e-05		
## 31	0.6370008 63 1.817114e-05		
## 36	0.6328245 63 2.397306e-05		
## 43	0.6270161 63 3.500542e-05		
## 151	0.6041192 63 1.447163e-04		

```
## 160 0.6024866 63 1.594511e-04
## 211 0.6016949 59 4.238812e-04
## 46 0.6006144 63 1.780847e-04
## 229 0.5914085 59 7.412314e-04
## 215 0.5402104 59 9.165006e-03
## 41 0.5322581 63 6.503995e-03
## 27 0.5307700 63 6.974040e-03
## 52 0.5302899 63 7.132298e-03
## 216 0.5276447 59 1.598323e-02
## 338 0.5264876 63 8.509354e-03
## 230 0.5053256 59 4.071285e-02
## 165 0.4987787 62 3.365520e-02
## 318 0.4970238 63 3.117452e-02
## 150 0.4954397 63 3.331712e-02
## 171 -0.5155530 63 1.397320e-02
## 173 -0.5209789 63 1.094818e-02
## 23 -0.5540515 63 2.251219e-03
## 210 -0.5689882 63 1.040665e-03
## 148 -0.5789076 62 7.542336e-04
## 228 -0.6133255 59 2.199641e-04
## 322 -0.6170076 59 1.777250e-04
## 224 -0.6760959 59 3.829287e-06
## 225 -0.7443519 59 1.285176e-08
## 103 -0.8238686 62 2.025047e-13
## 312 -0.8500667 62 0.000000e+00
## 90 -0.8714714 62 0.000000e+00
## NA NA NA NA
```

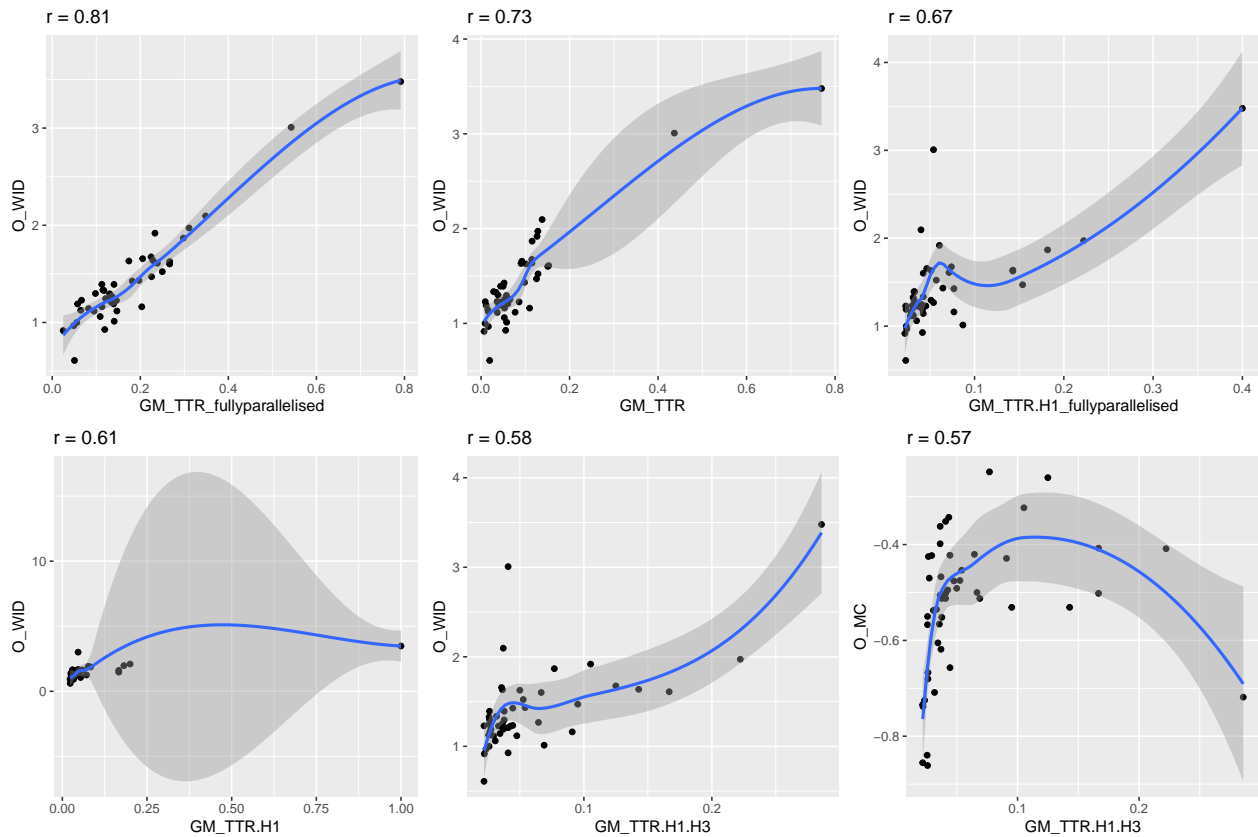
Positive Correlations

We here plot the six highest *positive* correlations (in terms of Spearman coefficients) which are still significant after the Bonferroni correction *and* which are found between measures proposed by *different participants* (there are many measures by the same participants that highly correlate). These are hand-picked from the lists above.

TRACK A

Plot the six significant correlations with highest Spearman coefficients for Track A. Warning messages are disabled since there are several NAs that throw errors.

```
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
```

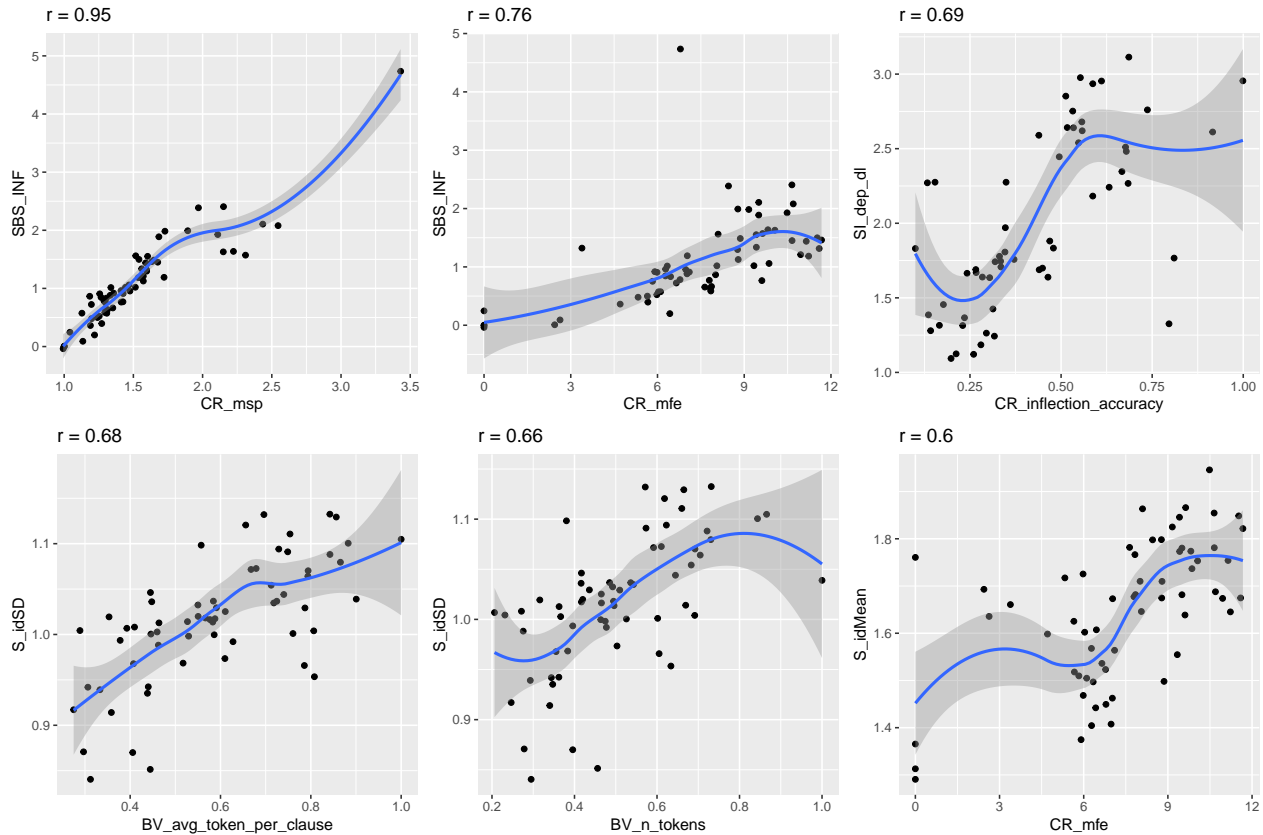


`ggsave("~/Github/ComplexityMetaAnalyses/Figures/Corrs/TrackA/track_a_plot_corrected.pdf", track.a.plot.)`

TRACK B

Plot the six significant correlations with highest Spearman coefficients for Track B. Warning messages are disabled since there are several NAs that throw errors.

```
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
```



`ggsave("~/Github/ComplexityMetaAnalyses/Figures/Corrs/TrackB/track_b_plot_corrected.pdf", track.b.plot.)`

Detailed Plots

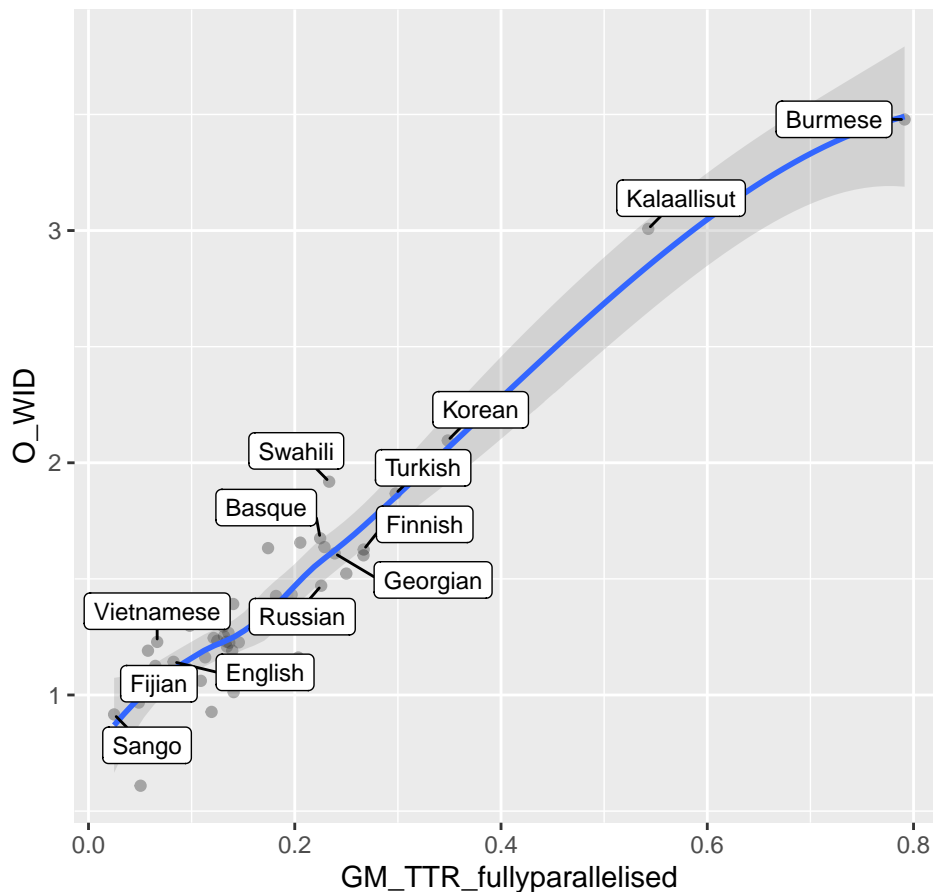
The code below adds labels to the points of plots, which helps with the interpretation of results. We here choose the two plots of Track A and Track B with the highest positive Spearman correlations.

TRACK A

```
#track.a <- track.a[track.a$id != "mya", ] # remove the outlier Burmese (mya)

track.a.plot1.detailed <- ggplot(track.a, aes(x = GM_TTR_fullyparallelised, y = O_WID)) +
  geom_point(alpha = 0.3) +
  geom_smooth(method = loess, alpha = 0.3) +
  geom_label_repel(data = track.a[track.a$language == "Fijian" | track.a$language == "Sango" | track.a$language == "Burmese"],
    min.segment.length = 0,
    #nudge_x = 0.1,
    aes(label = language),
    size = 3) +
  labs(title = paste("r = ",
    round(track.a.df[track.a.df$row == "GM_TTR_fullyparallelised" & track.a.df$col == "O_WID"], 2),
    sep = "")) +
  theme(legend.position = "none")
track.a.plot1.detailed
```

```
## `geom_smooth()` using formula 'y ~ x'
r = 0.81
```



```
ggsave("~/Github/ComplexityMetaAnalyses/Figures/Corrs/TrackA/track_a_plot1_detailed.pdf", track.a.plot1)
```

```
## `geom_smooth()` using formula 'y ~ x'
```

Some comments: This plot shows that the Type-Token Ratio (TTR) and the Word Information Density (WID) are highly correlated across the languages of the Parallel Bible Corpus sample. Burmese (mya) is an outlier here with very high TTR and WID. This is an artifact of the writing system, since it does not delimit orthographic words by white spaces, but rather phrases. For Kalaallisut, on the other hand, the result makes sense (if we accept the latinized writing proposed for this language). Some of the low TTR languages include Sango (sag), Fijian (fij), Thai (tha), and Yoruba (yor).

TRACK B

```
#track.b <- track.b[track.b$id != "uig", ] # remove the outlier Uyghur (uig)
```

```
track.b.plot1.detailed <- ggplot(track.b, aes(x = CR_msp, y = SBS_INF)) +
  geom_point(alpha = 0.3) +
  geom_smooth(method = loess, alpha = 0.3) +
  geom_label_repel(data = track.b[track.b$language == "Chinese" | track.b$language == "Vietnamese" | track.b$language == "Burmese"],
    min.segment.length = 0,
    #nudge_x = 0.1,
```




```

    aes(label = language),
    size = 3) +
  labs(title = paste("r = ",
    round(track.b.df[track.b.df$row == "CR_msp" & track.b.df$col == "SBS_INF", ]$corr,
    sep = "")) +
  theme(legend.position = "none")
track.b.plot1.detailed

```

```
## `geom_smooth()` using formula 'y ~ x'
```

$r = 0.95$



```
ggsave("~/Github/ComplexityMetaAnalyses/Figures/Corrs/TrackB/track_b_plot1_detailed.pdf", track.b.plot1
```

```
## `geom_smooth()` using formula 'y ~ x'
```

Some comments: This plot shows the correlation between the so-called Mean Size of Morphological Paradigms (MSP), which is defined by CR as “simply the number of word-form types divided by the number of lemma types”, and the difference in unigram entropy of word tokens in the original texts and the lemmatized texts (INF) as defined by SBS. It is certainly not unexpected, but reassuring, to see these measure highly correlated. The outlier to the high end Uyghur (uig) is likely *not* an artifact, as this language indeed has many productive morphological paradigms. Other languages to the high end of morphological complexity include Ancient Greek (gre), Classical Latin (lat), Turkish (tur), and Old Church Slavonic (chu). Languages to the low end are Vietnamese (vie), Indonesian (ind), Mandarin Chinese (cmn), and Afrikaans (afr). Note that the very low morphological complexity scores of Korean (kor) are an artifact of the way the Korean data is presented in the UD. Namely, the “lemmas” given for Korean are actually merely morphologically segmented forms rather than inflectionally neutralized forms as for the other languages. Thus, it makes sense that the MSP is

exactly 1 and the INF is 0.

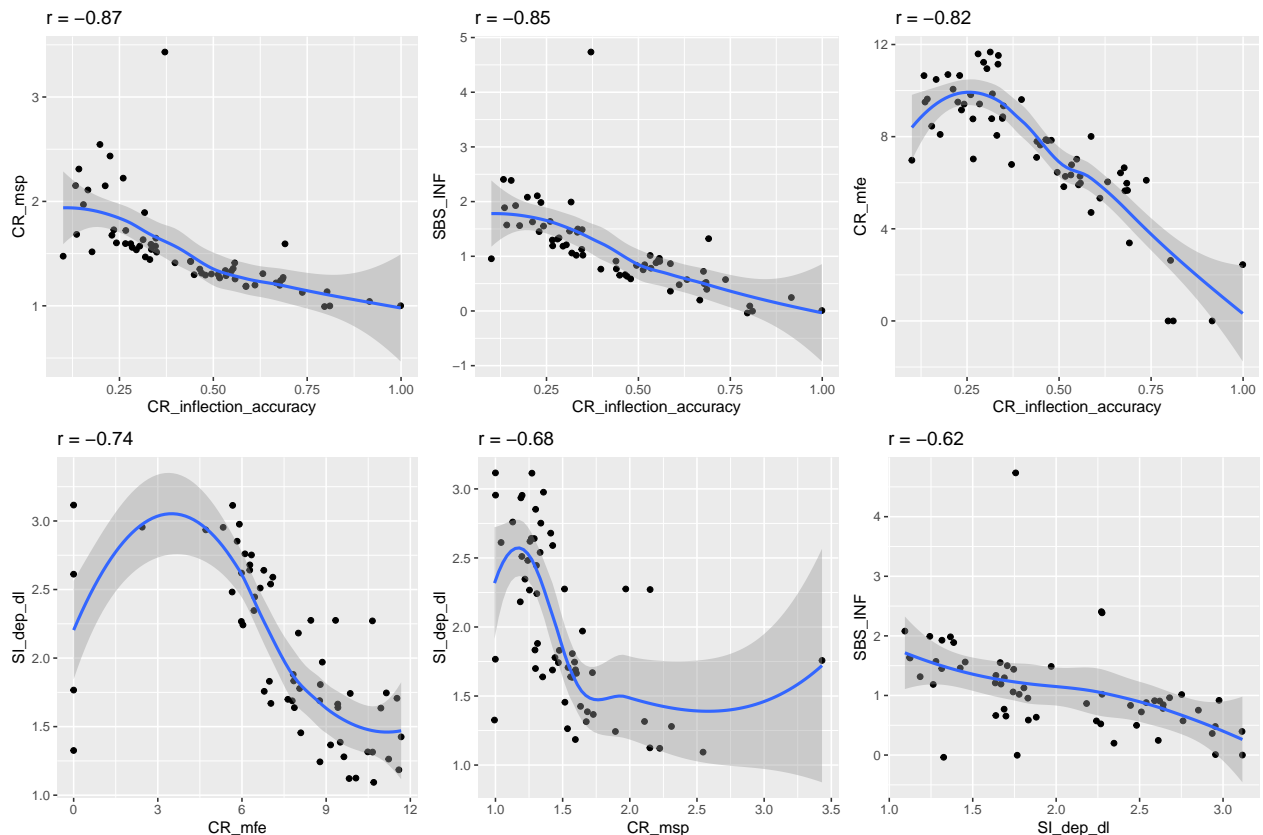
Negative Correlations

I here plot the six highest *negative* correlations (in terms of Spearman coefficients) which are still significant after the Bonferroni correction *and* which are found between measures proposed by *different participants* (there are many measures by the same participants that highly correlate). These are hand-picked from the lists above (could also be implemented more elegantly).

TRACK B (TRACK A does not yield negative correlations)

Plot the six significant correlations with lowest (i.e. negative) Spearman coefficients for Track B. Warning messages are disabled since there are several NAs that throw errors.

```
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
```




```
ggsave("~/Github/ComplexityMetaAnalyses/Figures/Corrs/TrackB/track_b_plot_negative_corrected.pdf", track_b_plot_negative_corrected)
```

Detailed Plots

```
#track.b <- track.b[track.b$id != "uyghur", ] # remove the outlier Uyghur (uyg)

track.b.plot1.negative.detailed <- ggplot(track.b, aes(x = CR_inflection_accuracy, y = SBS_INF)) +
  geom_point(alpha = 0.3) +
  geom_smooth(method = loess, alpha = 0.3) +
  geom_label_repel(min.segment.length = 0,
                  #nudge_x = 0.1,
                  aes(label = language),
                  size = 3) +
  labs(title = paste("r = ",
                    round(track.b.df[track.b.df$row == "CR_inflection_accuracy" & track.b.df$col == "SBS_INF",
                    sep = "")) +
  theme(legend.position = "none")
track.b.plot1.negative.detailed
```

```
## `geom_smooth()` using formula 'y ~ x'
```



```
ggsave("~/Github/ComplexityMetaAnalyses/Figures/Corrs/TrackB/track_b_plot1_negative_detailed.pdf", track.b.plot1.negative.detailed)
```


```
## `geom_smooth()` using formula 'y ~ x'
```

```
#track.b <- track.b[track.b$language != "Korean", ] # remove the outlier Korean
```

```
track.b.plot2.negative.detailed <- ggplot(track.b, aes(x = CR_mfe, y = SI_dep_dl)) +
  geom_point(alpha = 0.3) +
  geom_smooth(method = loess, alpha = 0.3) +
  geom_label_repel(data = track.b[track.b$language == "Chinese" | track.b$language == "Vietnamese" | track.b$language == "English" | track.b$language == "Ancient Greek" | track.b$language == "Greek" | track.b$language == "German" | track.b$language == "Russian" | track.b$language == "Uyghur" | track.b$language == "Finnish" | track.b$language == "Basque" | track.b$language == "Latin" | track.b$language == "Old Church Slavonic" | track.b$language == "Korean"],
    min.segment.length = 0,
    #nudge_x = 0.1,
    aes(label = language),
    size = 3) +
  labs(title = paste("r = ",
    round(track.b.df[track.b.df$row == "CR_mfe" & track.b.df$col == "SI_dep_dl", ]$correlation, 2),
    sep = "")) +
  theme(legend.position = "none")
track.b.plot2.negative.detailed
```

```
## `geom_smooth()` using formula 'y ~ x'
```

$r = -0.74$



```
ggsave("~/Github/ComplexityMetaAnalyses/Figures/Corrs/TrackB/track_b_plot2_negative_detailed.pdf", track.b.plot2.negative.detailed)
```

```
## `geom_smooth()` using formula 'y ~ x'
```

Conclusions

Some more general observations based on these analyses include:

- Many of the measures proposed by the same participants highly correlate. This is the case, for instance, for the measures proposed by GM in Track A, but also measures of BV and SI in Track B. In the case of GM, this is because many of the measures are virtually the same, but with minor shades of modification. In the case of BV, while at first sight the measures seem to conceptually differ, they essentially boil down to the same underlying causes. For example, the number of tokens in a sentence highly predicts the average maximal depth of a tree over the sentence. In the case of SI, the highly correlating measures in fact only have very few data points (only two or three in some cases). So, arguably all of these intra-participant correlations are somewhat artificial in the sense that they are either redundant or driven by inappropriate sample size.
- There are several strong positive correlations between simple measures relating to the number of types and tokens (GM_TTR_fullyparallelised, BV_n_tokens, etc.), and measures of information density (O_WID, S_idSD). Interestingly, this is the case for both tracks, since Oh used the Bible texts, and Semenuks used the UD. Information density is generally assumed to be a measure of “syntax” that has psycholinguistic relevance in terms of language processing. However, the fact that it is highly predictable by some of the simplest word frequency measures potentially goes to show that the underlying reasons for complexity are ironically quite simple.
- We have mainly discussed positive correlations here, meaning that certain measures are essentially (better or worse) replacements of other measures. In fact, the majority of correlations still significant after the Bonferroni correction are positive (in Track A all of them are positive, in Track B 37 out of 49 are positive). Some of the negative correlations we do find are between CR’s “inflection accuracy” and different measures of inflectional complexity (CR_mfe, SBS_INF, CR_msp). This makes perfect sense given that inflection accuracy is a measure that reflects the difficulty of NLP tools to automatically deal with inflectional morphology. The more complex the morphology, the lower the accuracy of the automated tool. A negative correlation that seems both robust and potentially interesting is that the dependency lengths in noun phrases with marked possessives (SI_dep_dl) apparently are in a clear trade-off with different measures of inflectional complexity. However, the fact that there are few such instances of robust negative correlations between measures of different domains suggests that there are relatively few clear complexity trade-offs. This is sth. to further think about.